

Celery on Docker Swarm

Wei Lin

Abstract

In the territory of Python, Celery is a famous distributed task queue framework, its Canvas mechanism is specially powerful in building complex workflow.

Celery can function in distributed environment and go perfectly along with Docker-Swarm: Docker-Swarm provides a cluster environment and the worker containers to sustain Celery; worker containers and processes and be dynamically scaled and expanded to fulfill Celery's need, work together parallelly to accomplish the computation.

With Docker-Swarm , a cluster will be built upon two Raspberry Pi machines. Hadoop entry-level "Word Count" program will be re-written in Python and executed parallelly via Celery on the cluster.

Steps:

1. Establish a Docker Swarm:

With reference to the article "[Let Docker Swarm all over your Raspberry Pi Cluster](http://blog.hypriot.com/post/let-docker-swarm-all-over-your-raspberry-pi-cluster/)" (<http://blog.hypriot.com/post/let-docker-swarm-all-over-your-raspberry-pi-cluster/>), a Docker-Swarm had been built upon two Raspberry Pi.

The swarm is composited fo two Docker machines:

- host rpi202(192.168.0.114) as Swarm Manager , the Docker machine name is "master01".
- host rpi201(192.168.0.109) as Swarm Node , the Docker machine name is "node01".

In []:

```
HypriotOS: pi@rpi202 in ~
$ docker-machine ls
NAME          ACTIVE    DRIVER    STATE    URL                                     SWARM
master01      True     hypriot   Running  tcp://192.168.0.114:2376             master01 (master)
node01        True     hypriot   Running  tcp://192.168.0.109:2376             master01
HypriotOS: pi@rpi202 in ~
$

# Nodes in the swarm:

HypriotOS: pi@rpi202 in ~
$ docker $(docker-machine config --swarm master01) info
Containers: 4
Images: 51
Role: primary
Strategy: spread
Filters: health, port, dependency, affinity, constraint
Nodes: 2
  master01: 192.168.0.114:2376
    L Status: Healthy
    L Containers: 3
    L Reserved CPUs: 0 / 4
    L Reserved Memory: 0 B / 972 MiB
    L Labels: executiondriver=native-0.2, kernelversion=4.1.8-hypriotos-v7+, operatings
  node01: 192.168.0.109:2376
    L Status: Healthy
    L Containers: 1
    L Reserved CPUs: 0 / 4
    L Reserved Memory: 0 B / 972 MiB
    L Labels: executiondriver=native-0.2, kernelversion=4.1.8-hypriotos-v7+, operatings
CPUs: 8
Total Memory: 1.899 GiB
Name: b7def5d9af98
HypriotOS: pi@rpi202 in ~
$
```

2. Establish the broker for Celery, with a Docker container running Redis.

In []:

```
HypriotOS: pi@rpi202 in ~
$ docker run -d -p 6379:6379 --net=mynet --name=redis --volume=/data:/data hypriot/rp
a2abf9277b5e4818da89ffa282a706506ef288426486cc25b431208564bf6e0f
```

```
HypriotOS: pi@rpi202 in ~
$ docker ps
CONTAINER ID          IMAGE                    COMMAND                  CREATED
a2abf9277b5e         hypriot/rpi-redis       "/entrypoint.sh redis"  13 hours ago
f0ce33ca1152         hypriot/rpi-swarm       "/swarm join --advert"  6 days ago
b7def5d9af98         hypriot/rpi-swarm       "/swarm manage --tlsv"  6 days ago
ad594813f8f0         nimblestratus/rpi-consul "/bin/start -server -"  6 days ago
HypriotOS: pi@rpi202 in ~
$
```

3. Copy `celeryconfig.py` , `start_workers.sh` , and the folder "`word_count`" to two hosts, under the folder of `/data/celery_projects`.

In []:

```
HypriotOS: pi@rpi202 in /data/celery_projects
$ ll
total 20
drwxr-xr-x 3  999 root 4096 Jan 25 23:01 ./
drwxr-xr-x 3  999 root 4096 Jan 25 23:01 ../
-rw-r--r-- 1  999 root 1079 Jan 25 21:12 celeryconfig.py
-rwxr-xr-x 1  999 root  732 Jan 25 22:53 start_workers.sh* <--- script to start up w
drwxr-xr-x 3 root root 4096 Jan 25 23:01 word_count/
HypriotOS: pi@rpi202 in /data/celery_projects
$
```

```
HypriotOS: pi@rpi201 in /data/celery_projects
$ ll
total 20
drwxr-xr-x 3 root root 4096 Jan 25 23:03 ./
drwxr-xr-x 3  999 root 4096 Jan 25 22:55 ../
-rw-r--r-- 1 root root 1079 Jan 25 21:12 celeryconfig.py
drwxr-xr-x 3 root root 4096 Jan 25 23:03 word_count/
HypriotOS: pi@rpi201 in /data/celery_projects
$
```

4. Establish Flower

(<http://docs.celeryproject.org/en/latest/userguide/monitoring.html#flower-real-time-celery-web-monitor>) container for monitoring.

In []:

```
HypriotOS: pi@rpi202 in ~
$ docker run -d -p 5555:5555 --net=mynet --name=flower --volume=/data/celery_projects
276f00591fd7042139ddf660730d223bcf19e9f8bd369f075de417140b6dfd4a
HypriotOS: pi@rpi202 in ~
$
```

```
HypriotOS: pi@rpi202 in ~
$ docker ps
CONTAINER ID          IMAGE                COMMAND              CREATED
276f00591fd7         wei1234c/celery_armv7  "/bin/sh -c 'cd /cele"  31 seconds ago
```

Connect to container of Flower, view with a web browser, no tasks yet.



5. Deploy Celery worker containers through Swarm-Manager.

deploy worker containers with :

start_workers.sh

In []:

```
# ./start_workers.sh

echo "Starting Celery cluster containers"

eval $(docker-machine env --swarm master01)

PROJECT=$1 # project name
WORKER_START_ID=$2 # worker container index, the first one.
WORKER_LAST_ID=$3 # worker container index, the last one.
CONCURRENCY=$4 # number of subprocesses per worker.

for (( i=${WORKER_START_ID}; i<=${WORKER_LAST_ID}; i=i+1 ))
do
    docker run -d --name=${PROJECT}_celery${i} --hostname=${PROJECT}_celery${i} --net=mynet
done
```

Establish Celery worker containers for the first round.

Four containers were established , each container has one Celery worker , each worker can have 5 subprocesses at most .

In []:

```
# CLI parameters:
# $1 # project name
# $2 # worker container index, the first one.
# $3 # worker container index, the last one.
# $4 # number of subprocesses per worker.

HypriotOS: pi@rpi202 in /data/celery_projects
$ ./start_workers.sh word_count 1 4 5
Starting Celery cluster containers _____
a22b08a0818b3246f90511ad21cb2a0ab37a4e72661bf559ade7e320db030505
77eabdded27e4ea3aaa640480c088fa7b4b9818fc3e40fb66636cc9abe8a78e69
df05a7204f40470cfd8eee21a06be45f5a306ea32df0196f3d004beac5d2f82d
e67d39740ace5c2a5b9a05e6ca1adc73c5e5944e62302d02391d37f7ee6aa479

# Four containers were established , each container has one Celery worker , each worker

HypriotOS: pi@rpi202 in /data/celery_projects
$ docker ps
CONTAINER ID          IMAGE                  COMMAND                  CREATED
e67d39740ace          wei1234c/celery_armv7  "/bin/sh -c 'cd /cele"  About a minut
df05a7204f40          wei1234c/celery_armv7  "/bin/sh -c 'cd /cele"  About a minut
77eabdded27e4          wei1234c/celery_armv7  "/bin/sh -c 'cd /cele"  About a minut
a22b08a0818b          wei1234c/celery_armv7  "/bin/sh -c 'cd /cele"  About a minut
276f00591fd7          wei1234c/celery_armv7  "/bin/sh -c 'cd /cele"  37 minutes ag
a2abf9277b5e          hypriot/rpi-redis      "/entrypoint.sh redis"   13 hours ago
980161d10fc4          hypriot/rpi-swarm      "/swarm join --advert"   6 days ago
f0ce33ca1152          hypriot/rpi-swarm      "/swarm join --advert"   6 days ago
b7def5d9af98          hypriot/rpi-swarm      "/swarm manage --tlsv"   6 days ago
ad594813f8f0          nimblestratus/rpi-consul  "/bin/start -server -"   6 days ago
HypriotOS: pi@rpi202 in /data/celery_projects
$
```

Flower shows the status. There are four workers.

- celery@worker1.word_count_celery1
- celery@worker2.word_count_celery2
- celery@worker3.word_count_celery3
- celery@worker4.word_count_celery4

In []:

```
# CLI parameters:
# $1 # project name
# $2 # worker container index, the first one.
# $3 # worker container index, the last one.
# $4 # number of subprocesses per worker.
```

HypriotOS: pi@rpi202 in /data/celery_projects

```
$ ./start_workers.sh word_count 5 8 5
```

Starting Celery cluster containers

```
a4de4967fd6211266cbad04fecfc357aa81789063cca3042388019adab2a6c71
7066ba6021de870f1332858c6f96673a159d7e5031a5682d3853fa6bd8fe2252
79792c823fbf2769e4983c525598c30ba3758c23697ef66a78a54163374d3233
25c02d07ed6f1217ee68dc486a6586262ca2e3ed01a2a8690eaa2a783ad07d73
```

HypriotOS: pi@rpi202 in /data/celery_projects

```
$ docker ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED
25c02d07ed6f	wei1234c/celery_armv7	"/bin/sh -c 'cd /cele"	18 seconds ag
79792c823fbf	wei1234c/celery_armv7	"/bin/sh -c 'cd /cele"	39 seconds ag
7066ba6021de	wei1234c/celery_armv7	"/bin/sh -c 'cd /cele"	About a minut
a4de4967fd62	wei1234c/celery_armv7	"/bin/sh -c 'cd /cele"	About a minut
e67d39740ace	wei1234c/celery_armv7	"/bin/sh -c 'cd /cele"	15 minutes ag
df05a7204f40	wei1234c/celery_armv7	"/bin/sh -c 'cd /cele"	15 minutes ag
77eabded27e4	wei1234c/celery_armv7	"/bin/sh -c 'cd /cele"	15 minutes ag
a22b08a0818b	wei1234c/celery_armv7	"/bin/sh -c 'cd /cele"	16 minutes ag
276f00591fd7	wei1234c/celery_armv7	"/bin/sh -c 'cd /cele"	51 minutes ag
a2abf9277b5e	hypriot/rpi-redis	"/entrypoint.sh redis"	14 hours ago
980161d10fc4	hypriot/rpi-swarm	"/swarm join --advert"	6 days ago
f0ce33ca1152	hypriot/rpi-swarm	"/swarm join --advert"	6 days ago
b7def5d9af98	hypriot/rpi-swarm	"/swarm manage --tlsv"	6 days ago
ad594813f8f0	nimblestratus/rpi-consul	"/bin/start -server -"	6 days ago

HypriotOS: pi@rpi202 in /data/celery_projects

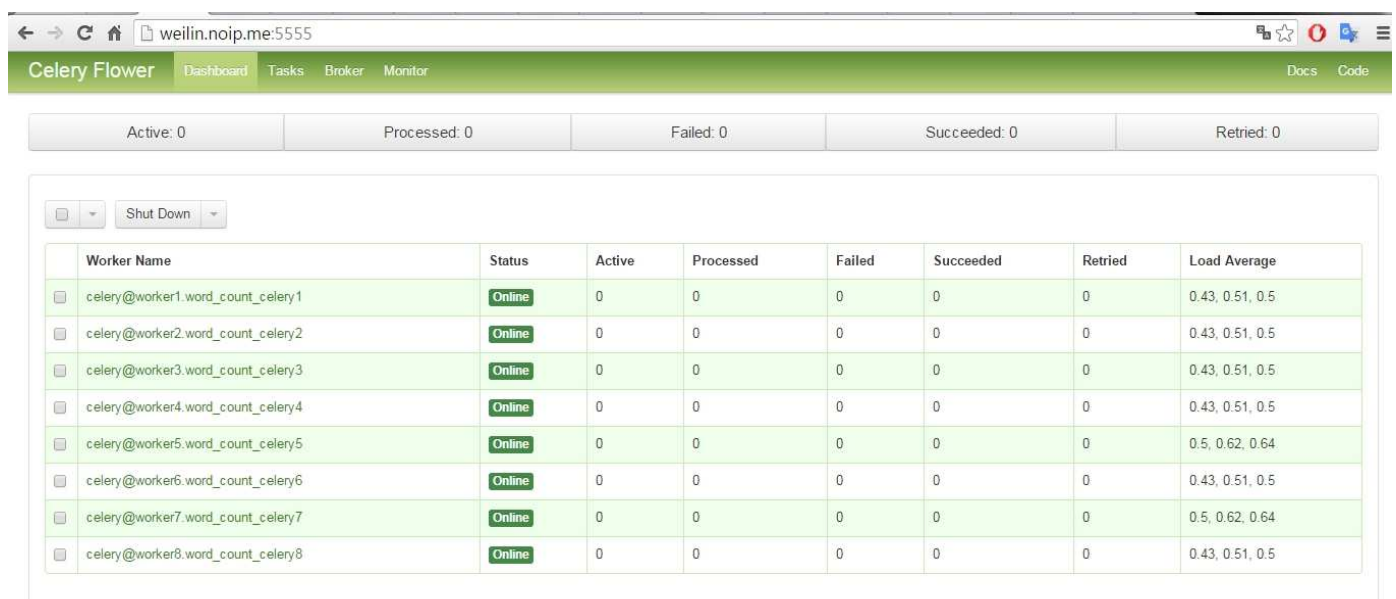
```
$
```

We can also view the list of Swarm nodes.

In []:

```
HypriotOS: pi@rpi201 in /data/celery_projects
$ docker ps
CONTAINER ID        IMAGE               COMMAND             CREATED
25c02d07ed6f        wei1234c/celery_armv7  "/bin/sh -c 'cd /cele"  16 minutes ago
7066ba6021de        wei1234c/celery_armv7  "/bin/sh -c 'cd /cele"  17 minutes ago
e67d39740ace        wei1234c/celery_armv7  "/bin/sh -c 'cd /cele"  32 minutes ago
df05a7204f40        wei1234c/celery_armv7  "/bin/sh -c 'cd /cele"  32 minutes ago
77eabdded27e4        wei1234c/celery_armv7  "/bin/sh -c 'cd /cele"  32 minutes ago
a22b08a0818b        wei1234c/celery_armv7  "/bin/sh -c 'cd /cele"  32 minutes ago
980161d10fc4        hypriot/rpi-swarm     "/swarm join --advert"   6 days ago
HypriotOS: pi@rpi201 in /data/celery_projects
$
```

Flower shows the workers. We have 8 of them now.



The screenshot shows the Celery Flower web interface in a browser. The top navigation bar includes 'Celery Flower', 'Dashboard', 'Tasks', 'Broker', and 'Monitor'. Below this, a summary row shows: Active: 0, Processed: 0, Failed: 0, Succeeded: 0, and Retried: 0. The main content area features a table with 8 workers, all in 'Online' status. Each worker row includes a checkbox, worker name, status, and various performance metrics.

	Worker Name	Status	Active	Processed	Failed	Succeeded	Retried	Load Average
<input type="checkbox"/>	celery@worker1.word_count_celery1	Online	0	0	0	0	0	0.43, 0.51, 0.5
<input type="checkbox"/>	celery@worker2.word_count_celery2	Online	0	0	0	0	0	0.43, 0.51, 0.5
<input type="checkbox"/>	celery@worker3.word_count_celery3	Online	0	0	0	0	0	0.43, 0.51, 0.5
<input type="checkbox"/>	celery@worker4.word_count_celery4	Online	0	0	0	0	0	0.43, 0.51, 0.5
<input type="checkbox"/>	celery@worker5.word_count_celery5	Online	0	0	0	0	0	0.5, 0.62, 0.64
<input type="checkbox"/>	celery@worker6.word_count_celery6	Online	0	0	0	0	0	0.43, 0.51, 0.5
<input type="checkbox"/>	celery@worker7.word_count_celery7	Online	0	0	0	0	0	0.5, 0.62, 0.64
<input type="checkbox"/>	celery@worker8.word_count_celery8	Online	0	0	0	0	0	0.43, 0.51, 0.5

"Word Count" program re-written in Python.

In [18]:

```
from word_count.tasks import *

# split text file into a list of words
def getWordsFromText(file = '.\\text\\test.txt'):
    with open(file) as f:
        lines = f.readlines()
    return ' '.join(lines).replace(',', ' ').replace('.', ' ').split()

def reduce(word_counts):
    wordCounts = {}

    for word_count in word_counts:
        if word_count is not None:
            wordCounts[word_count[0]] = wordCounts.get(word_count[0], 0) + word_count[1]

    result = sorted(list(wordCounts.items()),
                     key = lambda x: (x[1], x[0]),
                     reverse = True)

    return result
```

In [19]:

```
# list of words
words = getWordsFromText()
words[:3]
```

Out[19]:

```
["Aesop's", 'Fables', 'Translated']
```

In [20]:

```
# how many words in the article.
len(words)
```

Out[20]:

```
2190
```

./word_count/tasks.py , which contains mapper function:

```
from celery import group
from word_count.celery import app

@app.task
def mapper(word):
    return (word, 1) if len(word) >= 5 else None    # discard words which are too short.
```

Send messages to Celery workers, computing in parallel.

In [21]:

```
def count_celery(words):  
  
    # send messages to Celery workers.  
    asyncResults = [mapper.s(word).delay() for word in words] # mapper function is a  
    results = [asyncResult.get() for asyncResult in asyncResults if asyncResult.get()  
  
    return reduce(results)  
  
%time counts = count_celery(words)  
counts[:5]
```

Wall time: 3min 23s

Out[21]:

```
[('would', 12), ('which', 8), ('their', 8), ('caught', 6), ('Farmer',  
6)]
```

As messages sent , Flower shows the status.



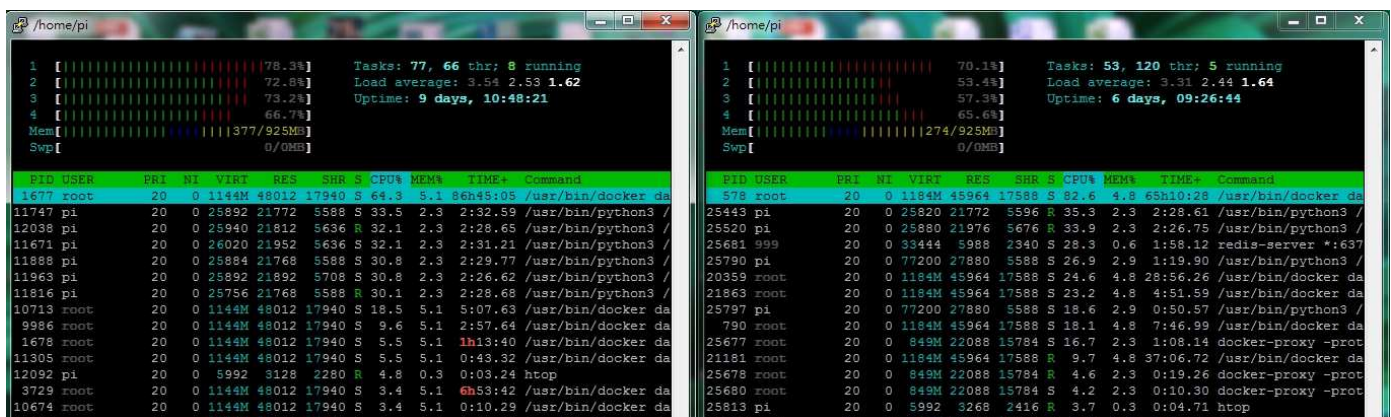
The screenshot shows the Celery Flower web interface in a browser. The address bar shows 'weilin.noip.me:5555'. The interface has a green header with 'Celery Flower' and navigation links: 'Dashboard', 'Tasks', 'Broker', and 'Monitor'. On the right of the header are links for 'Docs' and 'Code'. Below the header is a summary bar with five boxes: 'Active: 0', 'Processed: 1428', 'Failed: 0', 'Succeeded: 1427', and 'Retried: 0'. The main content area has a 'Shut Down' button and a table of workers.

Worker Name	Status	Active	Processed	Failed	Succeeded	Retried	Load Average
celery@worker1.word_count_celery1	Online	0	178	0	178	0	3.04, 1.12, 0.49
celery@worker5.word_count_celery5	Online	0	179	0	179	0	1.21, 0.8, 0.66
celery@worker7.word_count_celery7	Online	0	179	0	179	0	1.21, 0.8, 0.66
celery@worker4.word_count_celery4	Online	0	178	0	178	0	3.04, 1.12, 0.49
celery@worker2.word_count_celery2	Online	0	179	0	179	0	3.03, 1.14, 0.51
celery@worker6.word_count_celery6	Online	0	178	0	178	0	3.04, 1.12, 0.49
celery@worker8.word_count_celery8	Online	0	179	0	178	0	3.03, 1.14, 0.51
celery@worker3.word_count_celery3	Online	0	178	0	178	0	3.03, 1.14, 0.51

Finished.



Showing CPUs' load during the processing.



Summary:

With respect to Hadoop ecosystem, we can actually build a parallel-computing cluster in minutes with Celery and Docker-Swarm. This practice could be used as a handy tool in some circumstances.