

One for all, All for one (<https://www.youtube.com/watch?v=FcMCd520Ae8>)

SETI@home (<https://en.wikipedia.org/wiki/SETI@home>) liked Volunteer Computing (https://en.wikipedia.org/wiki/Volunteer_computing) on Mixed-Platforms Cluster - using Celery and Docker

Wei Lin
20160512

Experiment procedure:

Define Dockerfile for RPi:

Celery Worker Dockerfile for ARM v7
image name: wei1234c/one_for_all_all_for_one_armv7

```
In [ ]: # one_for_all_all_for_one_armv7
# Celery Worker Dockerfile
# for ARM v7
# 20160512

FROM wei1234c/celery_armv7

MAINTAINER Wei Lin

USER root

RUN mkdir /celery_projects

WORKDIR /celery_projects

COPY . /celery_projects/

RUN chmod +x /celery_projects/start_workers.sh

USER pi

CMD ["/bin/sh", "/celery_projects/start_workers.sh"]
```

Build Docker image for RPi:

image name: wei1234c/one_for_all_all_for_one_armv7

```
In [ ]: HypriotOS: pi@rpi202 in ~
$ docker build -t wei1234c/one_for_all_all_for_one_armv7 /dockerfiles/ARMv7/one_for_all_all_for_one
Sending build context to Docker daemon 9.216 kB
Step 1 : FROM wei1234c/celery_armv7
---> 8939b7e5c928
Step 2 : MAINTAINER Wei Lin
---> Using cache
---> 186c6ea155e8
Step 3 : USER root
---> Using cache
---> b84c84193d65
Step 4 : RUN mkdir /celery_projects
---> Using cache
---> fff839303a93
Step 5 : WORKDIR /celery_projects
---> Using cache
---> 27938d9d6ae8
Step 6 : COPY . /celery_projects/
---> c120021c8dc0
Removing intermediate container 5b4421e5472b
Step 7 : RUN chmod +x /celery_projects/start_workers.sh
---> Running in f719c550c90d
---> efa549c140dd
Removing intermediate container f719c550c90d
Step 8 : USER pi
---> Running in 1c7a5587769b
---> 427f37517ba9
Removing intermediate container 1c7a5587769b
Step 9 : CMD /bin/sh /celery_projects/start_workers.sh
---> Running in ed24629cb0ba
---> d7bb6603b6c6
Removing intermediate container ed24629cb0ba
Successfully built d7bb6603b6c6
HypriotOS: pi@rpi202 in ~
$
```

Define Dockerfile for amd64:

Celery Worker Dockerfile for amd64

image name: wei1234c/one_for_all_for_one

```
In [ ]: # one_for_all_all_for_one
# Celery Worker Dockerfile
# for amd64
# 20160512

FROM ubuntu

MAINTAINER Wei Lin

USER root

# Add user pi
RUN \
    useradd -G adm,sudo,users -s /bin/bash -m pi && \
    echo 'pi:raspberry' | chpasswd

#RUN pip3 install pandas

# Install Python. _____
RUN apt-get update && \
    apt-get install -y python3 python3-pip python3-dev python3-numpy python3-scipy python3-matplotlib python3-pandas && \
    apt-get install -y python python-pip python-dev

# Install Celery _____
RUN \
    pip3 install -U celery

RUN \
    pip3 install -U redis

RUN mkdir /celery_projects

WORKDIR /celery_projects

COPY . /celery_projects/

RUN chmod +x /celery_projects/start_workers.sh

USER pi
```

```
CMD ["/bin/sh", "/celery_projects/start_workers.sh"]
```

Build Docker image for amd64:

image name: wei1234c/one_for_all_for_one

```
In [ ]: wei@Wei-Lenovo:~$ docker build -t wei1234c/one_for_all_all_for_one /docker/dockerfiles/amd64/one_for_all_all_for_one
Sending build context to Docker daemon 9.728 kB
Step 1 : FROM ubuntu
---> c5f1cf30c96b
Step 2 : MAINTAINER Wei Lin
---> Using cache
---> 626cc4694d46
Step 3 : USER root
---> Using cache
---> 9212cedf802b
Step 4 : RUN useradd -G adm,sudo,users -s /bin/bash -m pi && echo 'pi:raspberry' | chpasswd
---> Using cache
---> a237ec2f3a84
Step 5 : RUN apt-get update && apt-get install -y python3 python3-pip python3-dev python3-numpy python3-scipy python3
---> Using cache
---> f65b7e004075
Step 6 : RUN pip3 install -U celery
---> Using cache
---> 3764134da5f1
Step 7 : RUN pip3 install -U redis
---> Using cache
---> 61772207fc08
Step 8 : RUN mkdir /celery_projects
---> Using cache
---> c68f9dc73b5c
Step 9 : WORKDIR /celery_projects
---> Using cache
---> b9e490c48b98
Step 10 : COPY . /celery_projects/
---> Using cache
---> 55e921f0a082
Step 11 : RUN chmod +x /celery_projects/start_workers.sh
---> Using cache
---> 194e82a97639
Step 12 : USER pi
---> Using cache
---> 740675730169
Step 13 : CMD /bin/sh /celery_projects/start_workers.sh
---> Using cache
---> d77e8341bf85
Successfully built d77e8341bf85
```

```
wei@Wei-Lenovo:~$
```

Run container for Celery Broker, using Redis

```
In [ ]: HypriotOS: pi@rpi202 in ~
$ docker run -d -p 6379:6379 --name=redis --volume=/data:/data hypriot/rpi-redis
2ee100973b0e1317e7511de0c97b2a29ad02a688f9928c14f347922a4aa3fb5d

HypriotOS: pi@rpi202 in ~
$ docker ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS
2ee100973b0e	hypriot/rpi-redis	"/entrypoint.sh redis"	8 seconds ago	Up 7 seconds	0.0.0.0:6379->6379

```
HypriotOS: pi@rpi202 in ~
$
```

Run Flower (<http://docs.celeryproject.org/en/latest/userguide/monitoring.html#flower-real-time-celery-web-monitor>) container for monitoring

```
In [ ]: HypriotOS: pi@rpi202 in ~
$ docker run -d -p 5555:5555 --name=flower wei1234c/one_for_all_all_for_one_armv7 /bin/sh -c "cd /celery_projects && celery
3c6e9e85417b536d07562575711e5f288097ed48d6f12c0129155d01ea746e66

HypriotOS: pi@rpi202 in ~
$ docker ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS
3c6e9e85417b	wei1234c/one_for_all_all_for_one_armv7	"/bin/sh -c 'cd /cele"	9 seconds ago	Up 7 seconds
2ee100973b0e	hypriot/rpi-redis	"/entrypoint.sh redis"	2 hours ago	Up 2 hours

```
HypriotOS: pi@rpi202 in ~
$
```

No worker running, nothing shows in Flower.

← → ↺ 🏠 192.168.0.114:5555

Celery Flower Dashboard Tasks Broker Monitor Docs Code

Active: 0 Processed: 0 Failed: 0 Succeeded: 0 Retried: 0

☐ Shut Down

Worker Name	Status	Active	Processed	Failed	Succeeded	Retried	Load Average
-------------	--------	--------	-----------	--------	-----------	---------	--------------



Run Celery worker container on Raspberry Pi

Volunteer can join the cluster for distributed parallel computing, simply by running this Docker image.

```
In [ ]: HypriotOS: pi@rpi202 in ~
$ docker run -d --name=musketeer1 wei1234c/one_for_all_all_for_one_armv7
5286a81ba451b8473ef2b5e3bb965a7b0fc31511e1ed4e368659aece2478e053

HypriotOS: pi@rpi202 in ~
$ docker ps
CONTAINER ID        IMAGE                                     COMMAND                  CREATED            STATUS
5286a81ba451        wei1234c/one_for_all_all_for_one_armv7  "/bin/sh /celery_proj"  4 seconds ago     Up 3 seconds
3c6e9e85417b        wei1234c/one_for_all_all_for_one_armv7  "/bin/sh -c 'cd /cele"  5 minutes ago     Up 5 minutes
2ee100973b0e        hypriot/rpi-redis                       "/entrypoint.sh redis"  2 hours ago       Up 2 hours
HypriotOS: pi@rpi202 in ~
$
```

After Celery worker container started, it showed up in Flower. However, no task message was received yet.

← → ↻ 🏠 192.168.0.114:5555     

Celery Flower Dashboard Tasks Broker Monitor Docs Code

Active: 0 Processed: 0 Failed: 0 Succeeded: 0 Retried: 0

☐ ▼ Shut Down ▼

	Worker Name	Status	Active	Processed	Failed	Succeeded	Retried	Load Average
<input type="checkbox"/>	celery@worker.5286a81ba451	Online	0	0	0	0	0	0.09, 0.13, 0.26

In package "stock" there is a file "tasks.py", containing function "get_table" with which we define task message.

```

In [ ]: from stock.celery import app

import pandas as pd
from datetime import datetime

def get_url(stock_id, year = datetime.today().year, month = datetime.today().month):
    return 'http://www.twse.com.tw/ch/trading/exchange/STOCK_DAY/genpage/Report{year}{month:02}/{year}{month:02}_F3_1_8_{

@app.task
def get_table(stock_id, year = datetime.today().year, month = datetime.today().month):

    url = get_url(stock_id, year, month)
    targetTableIndex = 0

    table = pd.read_html(url,
                        attrs = {'border': '0' ,
                                'width': '598',
                                'align': 'center',
                                'cellpadding': '0',
                                'cellspacing': '1',
                                'class': 'board_trad'},
                        header = 1
                        )[targetTableIndex]

    table['stock_id'] = stock_id
    table = table.reindex(columns = ['stock_id', 'date', 'quantity', 'amount', 'open', 'highest', 'lowest', 'close', 'off

    return table.tail(1).values

```

```

In [14]: # Load stock.tasks, which contains the definition of function "get_table".
from stock.tasks import *
import numpy as np

```

```
In [7]: # Excute get_table funtion directly from local host (OS: Windows 7)
# No task message was sent to Celery broker
get_table(2356)
```

```
Out[7]: array([[2356, '105/05/12', 5614182, 114086911, 20.2, 20.55, 20.1, 20.3,
-0.1, 3159]], dtype=object)
```

Asynchronous function call

In IPython Notebook on local host, we sent a task message to Celery cluster, demanding computing service. There is no Docker mechanism on local host (OS: Windows 7).

With "get_table.apply_async()" a task message will be sent to Celery Broker.

Celery Broker will put the task message into a queue.

Worker on Raspberry Pi will pick up the message from queue and excute it, and return the result.

```
In [8]: r = get_table.apply_async(args = [2356])
r.get()
```

```
Out[8]: array([[2356, '105/05/12', 5614182, 114086911, 20.2, 20.55, 20.1, 20.3,
-0.1, 3159]], dtype=object)
```

Soon after we sent a task message to Celery broker, it shows in Flower that there is a task processed successfully.

The screenshot shows the Celery Flower web interface in a browser. The address bar displays '192.168.0.114:5555'. The interface has a green header with 'Celery Flower' and navigation links for 'Dashboard', 'Tasks', 'Broker', and 'Monitor'. On the right of the header are links for 'Docs' and 'Code'. Below the header, a summary bar shows: 'Active: 0', 'Processed: 1', 'Failed: 0', 'Succeeded: 1', and 'Retried: 0'. The main content area features a 'Shut Down' button and a table with worker information.

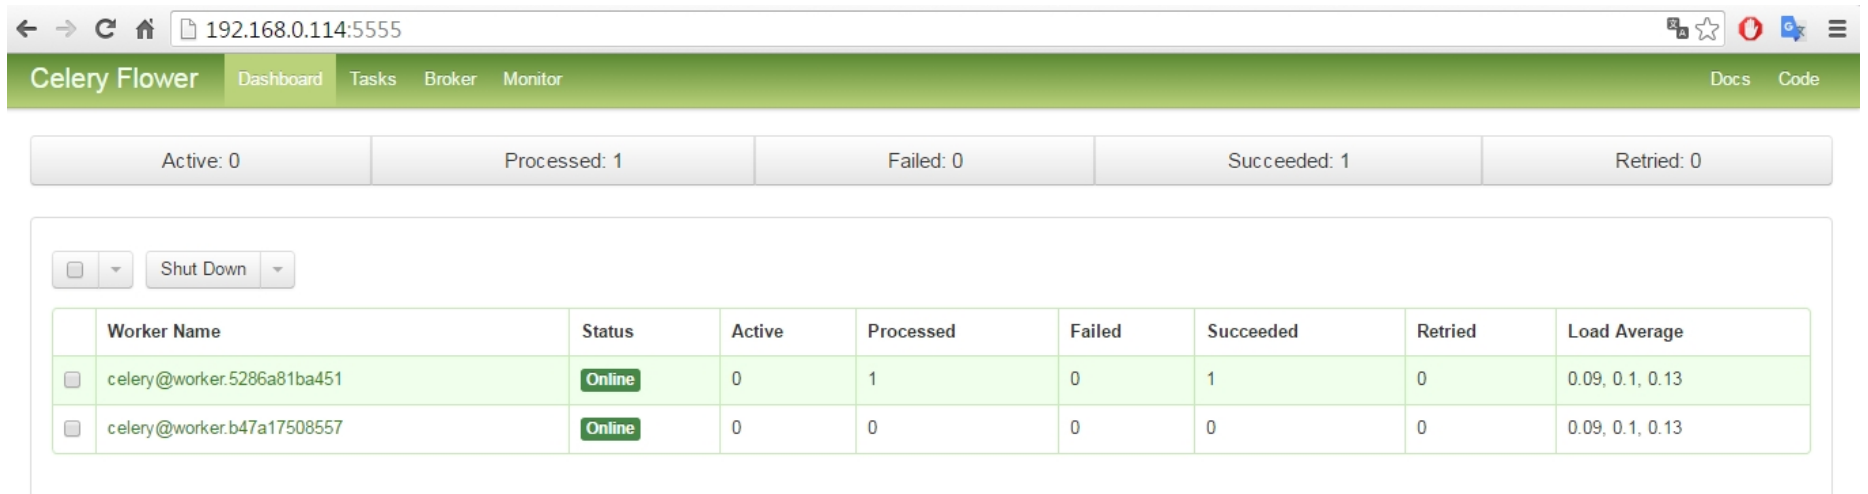
	Worker Name	Status	Active	Processed	Failed	Succeeded	Retried	Load Average
<input type="checkbox"/>	celery@worker.5286a81ba451	Online	0	1	0	1	0	0.08, 0.05, 0.12

In Raspberry Pi, we run another Celery worker container.

```
In [ ]: HypriotOS: pi@rpi202 in ~
$ docker run -d --name=musketeer2 wei1234c/one_for_all_all_for_one_armv7
b47a17508557cd48bbb21d48d7ad6b652e492058cf209d1c920a28db361e3568

HypriotOS: pi@rpi202 in ~
$ docker ps
CONTAINER ID        IMAGE                                     COMMAND                  CREATED            STATUS
b47a17508557        wei1234c/one_for_all_all_for_one_armv7  "/bin/sh /celery_proj"  4 seconds ago     Up 3 seconds
5286a81ba451        wei1234c/one_for_all_all_for_one_armv7  "/bin/sh /celery_proj"  21 minutes ago    Up 21 minutes
3c6e9e85417b        wei1234c/one_for_all_all_for_one_armv7  "/bin/sh -c 'cd /cele"  26 minutes ago    Up 26 minutes
2ee100973b0e        hypriot/rpi-redis                       "/entrypoint.sh redis"  3 hours ago       Up 3 hours
HypriotOS: pi@rpi202 in ~
$
```

We can see in Flower that we have two workers running now.



The screenshot shows the Celery Flower web interface in a browser. The address bar displays '192.168.0.114:5555'. The interface has a green header with 'Celery Flower' and navigation links for 'Dashboard', 'Tasks', 'Broker', and 'Monitor'. On the right of the header are links for 'Docs' and 'Code'. Below the header, a summary bar shows: 'Active: 0', 'Processed: 1', 'Failed: 0', 'Succeeded: 1', and 'Retried: 0'. The main content area features a 'Shut Down' button and a table of workers.

	Worker Name	Status	Active	Processed	Failed	Succeeded	Retried	Load Average
<input type="checkbox"/>	celery@worker.5286a81ba451	Online	0	1	0	1	0	0.09, 0.1, 0.13
<input type="checkbox"/>	celery@worker.b47a17508557	Online	0	0	0	0	0	0.09, 0.1, 0.13

On an AMD64 machine, we run two Celery worker containers.

```
In [ ]: wei@Wei-Lenovo:~$ docker run -d --name=musketeer3 wei1234c/one_for_all_all_for_one
65b6d885fb5bf06f10517c79325d19639446d939a0b2395aada323674e2eb121

wei@Wei-Lenovo:~$ docker run -d --name=musketeer4 wei1234c/one_for_all_all_for_one
d71f611c0ae6123f5139f02b1a5a5936f162259271061e7c256cc06d9b9d2511

wei@Wei-Lenovo:~$ docker ps
CONTAINER ID        IMAGE                                     COMMAND                  CREATED            STATUS
d71f611c0ae6        wei1234c/one_for_all_all_for_one        "/bin/sh /celery_proj"  6 seconds ago     Up 3 seconds
65b6d885fb5b        wei1234c/one_for_all_all_for_one        "/bin/sh /celery_proj"  About a minute ago Up About a minute
wei@Wei-Lenovo:~$
```

Now in Flower, we can see four workers - two workers on Raspberry Pi, and another two workers on the AMD64 machine.

← → ↺ 🏠

192.168.0.114:5555

🔖 🔴 🗨 ⋮

Celery Flower

Dashboard Tasks Broker Monitor

Docs Code

Active: 0

Processed: 1

Failed: 0

Succeeded: 1

Retried: 0

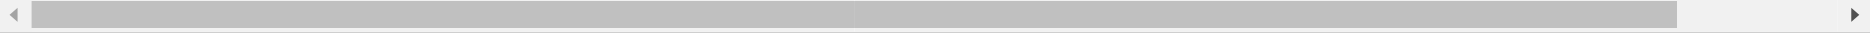
☐ ▾

Shut Down ▾

	Worker Name	Status	Active	Processed	Failed	Succeeded	Retried	Load Average
☐	celery@worker.65b6d885fb5b	Online	0	0	0	0	0	0.27, 0.57, 1.07
☐	celery@worker.b47a17508557	Online	0	0	0	0	0	0.05, 0.08, 0.05
☐	celery@worker.5286a81ba451	Online	0	1	0	1	0	0.05, 0.08, 0.05
☐	celery@worker.d71f611c0ae6	Online	0	0	0	0	0	0.27, 0.57, 1.07

```
In [29]: # The stock id list, about wich we will collect data.  
stocks = [1101, 1102, 1103, 1104, 1108, 1109, 1110, 1201, 1203, 1210, 1213, 1215, 1216, 1217, 1218, 1219, 1220,  
          1225, 1227, 1229, 1231, 1232, 1233, 1234, 1235, 1236, 1256, 1702, 1737, 1301]
```

```
In [30]: from pandas import DataFrame  
  
# collect data into a Pandas.DataFrame  
def reduce(results):  
    data = []  
    for result in results: data.append(result[0])  
  
    table = DataFrame(list(data),  
                      columns = ['stock_id', 'date', 'quantity', 'amount', 'open', 'highest', 'lowest', 'close', 'offset'])  
  
    return table
```



```
In [31]: def get_stock_prices(stocks):

    # send task messages to Celery broker
    asyncResults = [get_table.apply_async(args = [stock]) for stock in stocks]

    # get results from AsyncResults into a List
    results = [asyncResult.get() for asyncResult in asyncResults if asyncResult.get() is not None]

    return reduce(results)

%time prices = get_stock_prices(stocks)
prices[:5]
```

Wall time: 19.2 s

Out[31]:

	stock_id	date	quantity	amount	open	highest	lowest	close	offset	trades
0	1101	105/05/12	3242925	96161650	29.50	29.80	29.45	29.70	0.05	2064
1	1102	105/05/12	3094327	80204709	25.60	26.30	25.55	26.00	0.40	2046
2	1103	105/05/12	56511	492804	8.78	8.78	8.70	8.74	-0.04	32
3	1104	105/05/12	138794	2766347	20.15	20.15	19.85	20.00	-0.15	85
4	1108	105/05/12	85995	865446	10.05	10.10	10.00	10.05	0.00	43


```
In [32]: # list all results
prices
```

Out[32]:

	stock_id	date	quantity	amount	open	highest	lowest	close	offset	trades
0	1101	105/05/12	3242925	96161650	29.50	29.80	29.45	29.70	0.05	2064
1	1102	105/05/12	3094327	80204709	25.60	26.30	25.55	26.00	0.40	2046
2	1103	105/05/12	56511	492804	8.78	8.78	8.70	8.74	-0.04	32
3	1104	105/05/12	138794	2766347	20.15	20.15	19.85	20.00	-0.15	85
4	1108	105/05/12	85995	865446	10.05	10.10	10.00	10.05	0.00	43
5	1109	105/05/12	4000	40050	10.05	10.05	10.00	10.00	0.00	4
6	1110	105/05/12	31000	445500	14.20	14.45	14.20	14.45	0.00	13
7	1201	105/05/12	1015742	19595770	19.65	19.70	19.00	19.10	-0.60	384
8	1203	105/05/12	12004	257184	21.55	21.55	21.40	21.40	-0.30	11
9	1210	105/05/12	4576158	107022964	23.05	23.75	23.00	23.40	0.50	1946
10	1213	105/05/12	26010	447070	17.30	17.30	17.15	17.15	-0.10	21
11	1215	105/05/12	4658698	135637699	28.65	29.40	28.65	28.95	0.25	1992
12	1216	105/05/12	5823532	337920756	57.30	58.50	57.30	58.40	1.10	2945
13	1217	105/05/12	430450	3402141	7.85	8.00	7.81	7.85	0.00	173
14	1218	105/05/12	268733	3345762	12.45	12.55	12.35	12.40	-0.15	106
15	1219	105/05/12	37208	562998	15.20	15.20	15.00	15.05	-0.15	17
16	1220	105/05/12	50000	519600	10.45	10.50	10.30	10.45	0.00	25
17	1225	105/05/12	55382	1757142	31.65	31.90	31.20	31.20	-0.45	19
18	1227	105/05/12	511892	39680787	77.50	77.80	77.20	77.60	-0.10	401
19	1229	105/05/12	354427	7157504	20.40	20.40	20.05	20.10	0.10	183
20	1231	105/05/12	619161	18610131	29.25	30.55	29.25	30.20	1.05	328

	stock_id	date	quantity	amount	open	highest	lowest	close	offset	trades
21	1232	105/05/12	118000	9158100	77.30	77.90	77.30	77.40	0.10	55
22	1233	105/05/12	62200	2293460	36.80	37.05	36.80	36.85	0.05	32
23	1234	105/05/12	207595	6888635	33.60	33.60	32.95	32.95	-0.95	107
24	1235	105/05/12	25160	595842	23.45	23.95	23.30	23.90	0.00	23
25	1236	105/05/12	18010	361551	20.15	20.15	20.00	20.00	-0.15	16
26	1256	105/05/12	13150	1573200	120.00	120.00	118.50	118.50	-2.50	13
27	1702	105/05/12	1020054	63916873	61.70	63.20	61.70	62.60	0.60	601
28	1737	105/05/12	207912	5610143	27.00	27.10	26.85	26.85	-0.10	97
29	1301	105/05/12	3896203	300815138	77.00	77.40	76.80	77.10	0.10	2088

Tasks were distributed among four workers - two workers on Raspberry Pi, and another two workers on the AMD64 machine.

← → ↺ 🏠

192.168.0.114:5555

📄 ☆ 🔴 🗨️ ☰

Celery Flower

Dashboard Tasks Broker Monitor

Docs Code

Active: 0

Processed: 31

Failed: 0

Succeeded: 31

Retried: 0

☐

Shut Down

	Worker Name	Status	Active	Processed	Failed	Succeeded	Retried	Load Average
<input type="checkbox"/>	celery@worker.65b6d885fb5b	Online	0	9	0	9	0	0.01, 0.02, 0.16
<input type="checkbox"/>	celery@worker.b47a17508557	Online	0	6	0	6	0	0.36, 0.24, 0.17
<input type="checkbox"/>	celery@worker.5286a81ba451	Online	0	7	0	7	0	0.36, 0.24, 0.17
<input type="checkbox"/>	celery@worker.d71f611c0ae6	Online	0	9	0	9	0	0.01, 0.02, 0.16

